

1 Topics Review

1.1 Math Fundamentals

- **Recurrence Relations:**
 - Solve simple ones by writing out some terms and finding a pattern
 - Try to use Master's Theorem, or make a substitution first and then use it
 - Draw out a recursion tree, determine the number of levels and amount of work done at each layer.
- **Big-O Notation:** You should be very comfortable with the definitions for O , o , Ω , ω , and Θ .

1.2 Data Structures

- Arrays, Linked Lists, Stacks, and Queues.
- Heaps and Priority Queues
- Representations of a Graph – Adjacency Matrix and Adjacency List

1.3 Algorithms

You should be able to state how each of the following algorithms work and state its run-time.

- **Multiplication.** Karatsuba's Algorithm, Matrix Multiplication, Repeated Squaring
- **Mergesort.** Divide and Conquer, merging sorted sublists together.
- **Depth First Search.** Key idea is to use a stack to enumerate nodes. Assign **preorder** and **postorder** when pushing and popping nodes from the stack. Know what **tree edges**, **forward edges**, **back edges**, and **cross edges** are.
 - **Topological Sort.** Useful for ordering nodes in a Directed Acyclic Graph
 - **Strongly Connected Components.** Turn any graph into a DAG of SCCs.
- **Breadth First Search.** Key idea is to use a queue to store unprocessed nodes. Gives shortest paths in an unweighted graph.
- **Dijkstra's Algorithm** and **heaps.** Use DFS with a priority queue. Only works for positive edge weights.
- **Bellman Ford.** Recursive, takes advantage of the optimality of subpaths. Works on arbitrary graphs. Bonus: finds negative cycles!
- **Floyd-Warshall.** Finds all pairs shortest paths

- **Prim's Algorithm.** Build a tree from a single seed vertex. Uses the **Cut Property** to choose the smallest edge leaving the group.
- **Kruskal's Algorithm.** Build a tree by sorting the edges, joining disjoint groups of vertices. Uses the **Disjoint Set** data structure.

1.4 Algorithm Strategies

- **Greedy.** Pick the best option at each step. e.g., Prim's Algorithm or Kruskal's Algorithm
- **Divide and Conquer.** Split the problem into smaller subproblems (often in half), solve the subproblems, and combine the results. e.g., Mergesort, Integer Multiplication, Strassen's
- **Modify the Problem.** Build a graph from your problem or modify the given graph such that running an algorithm we talked about in class gives you the correct solution instantly.
- **Modify the Algorithm.** Make a slight modification to one an algorithm discussed in class to fit a particular problem.

2 Practice Problems

2.1 Asymptotic Notation

Exercise. Give a counterexample to the following statement. Then propose an easy fix.

If $f(n)$ and $g(n)$ are positive functions, then $f(n) + g(n) = \Theta(\min(f(n), g(n)))$

Solution.

This is false. Let $f(n) = n$, $g(n) = 1$. Then, $\min(f(n), g(n)) = 1$ but $f(n) + g(n) = n + 1$ is not $\Theta(\min(f(n), g(n))) = \Theta(1)$.

However, it is true that $f(n) + g(n) = \Theta(\max(f(n), g(n)))$. To prove this, note that:

$$\max(f(n), g(n)) \leq f(n) + g(n) \leq 2 \max(f(n), g(n)).$$

2.2 Number of Paths in a DAG

Exercise. Given a directed acyclic graph $G = (V, E)$ and two nodes $u, v \in V$, calculate the number of distinct paths from u to v . (Two paths are considered distinct if they have at least 1 vertex not in common)

Solution.

Since G is a DAG, we can perform a topological sort of its vertices in $O(n + m)$ time using DFS. Suppose our topological sort comes out to be $u, w_1, w_2, \dots, w_k, v$. Note that if v appears before u , then we should return 0 immediately since there would be no way to get from u to v .

Now, we define a recursive function $f(i)$ to be the number of distinct paths from u to w_i and let $u = w_0$. We are interested in the value $f(v)$.

Base Case: if $i = 0$, then $f(i) = 1$. There is only 1 way to get from u to u , which is to stay at u .

Recursive Case: if $i \neq 0$, then

$$f(i) = \sum_{w_j: (w_j, w_i) \in E} f(j)$$

The number of ways to get to w_i is the sum of the number of ways to get to each of w_i 's predecessors that have an edge connecting to w_i . Note that $0 \leq j < i$ and since the graph is topologically sorted, we guarantee that if there's an edge $w_j \rightarrow w_i$, then $j < i$.

Run-Time: First, the topological sort takes $O(n + m)$ time by DFS. Next, there can be up to n values of $f(i)$ to compute, and for each vertex w_j , it takes $O(1) + O(\text{indegree}(v))$ time to compute $f(w_j)$ so the total run-time is $O(n + m)$. The sum of the indegrees of each vertex in a directed graph is the number of edges in the graph. Note that we can pre-compute G^R in order to be able to access indegrees rather than just outdegrees in $O(1)$ time.

2.3 Semiconnectedness

Exercise. A directed graph $G = (V, E)$ is called *semiconnected* if for each pair of distinct vertices $u, v \in V$, there is either a path from u to v or a path from v to u . Find an algorithm to determine whether a directed graph is semiconnected.

Solution.

Let G' be the DAG on the strongly connected components of G . Number the nodes of G' in some topological order. We claim that G is semiconnected iff there is always a path from the i -th node to the j -th node of G' if $i < j$.

Suppose that there is always a path from the i -th node to the j -th node of G' if $i < j$. Then, for any two vertices u and v in G , they either belong to the same SCC (in which case there is a path from u to v and a path from v to u), or they belong to two different SCC's, c_u and c_v . Without loss of generality, suppose c_u occurs earlier than c_v with respect to the topological order. Then, by hypothesis, there is a path from c_u to c_v , so there is a path from u to v as desired.

Now suppose that for some topological sort of G' , there is no path from the i -th node (with respect to this sort) to the j -th node, where $i < j$. Then, there can be no path from the j -th node to the i -th node since $i < j$ and the nodes are topologically sorted. Hence, if u is a vertex lying in the i -th SCC and v a vertex lying in the j -th SCC, there are no paths from u to v or from v to u , so G is not semiconnected.

This observation allows us to devise an algorithm to determine if G is semiconnected: we first compute the strongly connected components of G and construct the graph G' in time $O(n + m)$, where $n = |V|$ and $m = |E|$. Then, using DFS from any vertex in G' , we can topologically sort G' in $O(n + m)$ time. Now, we check if there is a path from the i -th SCC to the j -th SCC if $i < j$. Note that this is the case iff there is an edge in G' from the i -th SCC to the $i + 1$ -th SCC. Hence, we can just scan through the topological sort in $O(n)$ time to see if the i -th SCC has an edge to the $i + 1$ -th SCC. The total running time of the algorithm is $O(n + m)$.

2.4 Negative Cycles

Exercise. *Modify Bellman Ford to set $\text{dist}[v] = -\infty$ for all vertices v that can be reached from the source via a negative cycle.*

Solution.

Normally when we check to see whether a graph has a negative cycle, we check to see that $f(n, v) = f(n - 1, v)$ for all $v \in V$. Since the graph has n vertices, any path that uses n edges must have touched some vertex twice because a path using n distinct vertices would be $n + 1$ edges long. Therefore, the path given by $f(n, v)$ could potentially have taken a cycle. If $f(n, v) < f(n - 1, v)$ for any $v \in V$, then we can say "Negative Cycle Detected".

It is important to see that if a negative cycle does exist in a graph, the condition that $f(n, v) = f(n - 1, v)$ will not be broken for *all* vertices that can be reached from the source via a negative cycle. For some vertices that are reachable by a negative cycle, n vertices are not enough to have traveled down a negative cycle.

To solve this problem, we compute $f(2n, v)$ for all $v \in V$. In up to $2n$ vertices, we can get from the source to any vertex v as well as have enough edges to travel along any negative cycle in the graph. If there is a negative cycle from the source to v , then the path itself without the negative cycle has up to n vertices, while the negative cycle has up to n vertices as well. Any negative cycle of more than n vertices can be decomposed into 2 smaller cycles, at least 1 of which is a negative cycle.

The run-time of this solution is still $O(nm)$ because we are simply doing $2nm$ work.


2.5 Minimum Spanning Tree True or False

Exercise. *True or False*


- *If G has a cycle with a unique heaviest edge e , then e cannot be part of any MST*
- *The shortest-path tree computed by Dijkstra's algorithm is necessarily an MST.*
- *Prim's algorithm works with negative weighted edges.*
- *If G has a cycle with a unique lightest edge e , then e must be part of every MST.*

Solution. • True. Suppose e was part of some spanning tree. If we remove e , we get two sets S (containing e) and $V - S$ with no edges in the MST going from S to $V - S$. Now, some other edge e' in the cycle must go from S to $V - S$, and replacing e with e' , we get another spanning tree. The total weight of this spanning tree is less than the weight of the former spanning tree since $w(e') < w(e)$, so the former spanning tree could not be minimal.

- False. For example, consider the following graph. The shortest path tree from A includes edges (A, B) , (A, D) , and (A, C) , but a minimal spanning tree (for example, (D, B) , (D, C) , and (D, A)) has total weight 6.



- True. The proof that Prim's algorithm is correct (the cut property) does not require edge weights to be negative. Another way to see this is that if we add a constant to all edges in the graph, the minimum spanning tree remains the same (since all spanning trees have total weight increased by the same amount), so adding a large enough constant, all edge weights are positive.
- False. Consider the following graph. The cycle BDE has unique smallest edge (B, D) , but is not in the spanning tree with edges (A, D) , (A, B) , (B, C) , (C, E) .



2.6 Minimum Spanning Tree with Few Edges

Exercise. Given a weighted graph with n vertices and $m \leq n + 10$ edges, show how to compute a minimum spanning tree in $O(n)$ time.

Solution.

First check if G is connected with a DFS ($O(m + n) = O(n)$ time). If G is not connected, there exists no MST, so return impossible.

Now suppose G is connected. Initially let the set T consist of all m edges. Let $c = m - (n - 1)$, i.e. let there be c too many edges in T . By the condition $m \leq n + 10$ we know that $c \leq 11$. Run the following

procedure c times:

DFS from an arbitrary vertex, keeping a predecessor array for the vertices. There exists some cycle, and thus there will be some back edge (u, v) . We know that u, v is a part of a cycle, so tracing back the predecessor array from u , we eventually reach v . Thus, we have detected a cycle in $O(n)$ time. Let $e = (u, v)$ be some edge in this cycle C with maximal weight (if there are ties choose e arbitrarily). Suppose there exists any tree T with edge e . We will show that there exists another tree T' without e with total weight less than or equal to T . Indeed, let $S \subset V$ be the set of vertices $s \in V$ such that in $T \setminus \{e\}$, s is connected to u (and hence not v). Then let the vertices of C be in the order u, x_1, \dots, x_k, v . Then some edge $e' = (x_i, x_{i+1})$ connects S to $V \setminus S$ and by the maximality of e we have $w(e) \geq w(e')$. Thus we can add e' , which creates a cycle

$$u \rightarrow v \rightarrow x_i \rightarrow x_{i+1} \rightarrow u$$

which can be broken by removing e . Since $w(e') \leq w(e)$ we have a new tree with total weight less than T .

Thus, there is an MST without e , so we can remove e and repeat the process on the new graph $G \setminus \{e\}$. After c repetitions, we will have a graph with $n - 1$ edges, and hence a tree. This will be an MST because we know that after each edge we remove, there exists an MST among the remaining edges (by the argument above).

This runs in $O(11(m + n)) = O(m + n) = O(n)$ time.